

Q1. When should you NOT use DynamoDB?

Q2. What causes a hot partition in DynamoDB?

Q3. How does DynamoDB scale automatically?

Q4. Difference between RDS and DynamoDB?

Q5. Explain how the COPY command works in Redshift.

Coding:

You have a table events with columns:  
event\_id, user\_id, event\_type, event\_time.

**Task:**

Write a SQL query to get the **latest event for each user**.

**Follow-up (Redshift):**

Why can this query be expensive in Redshift, and how would you optimize it for large datasets?